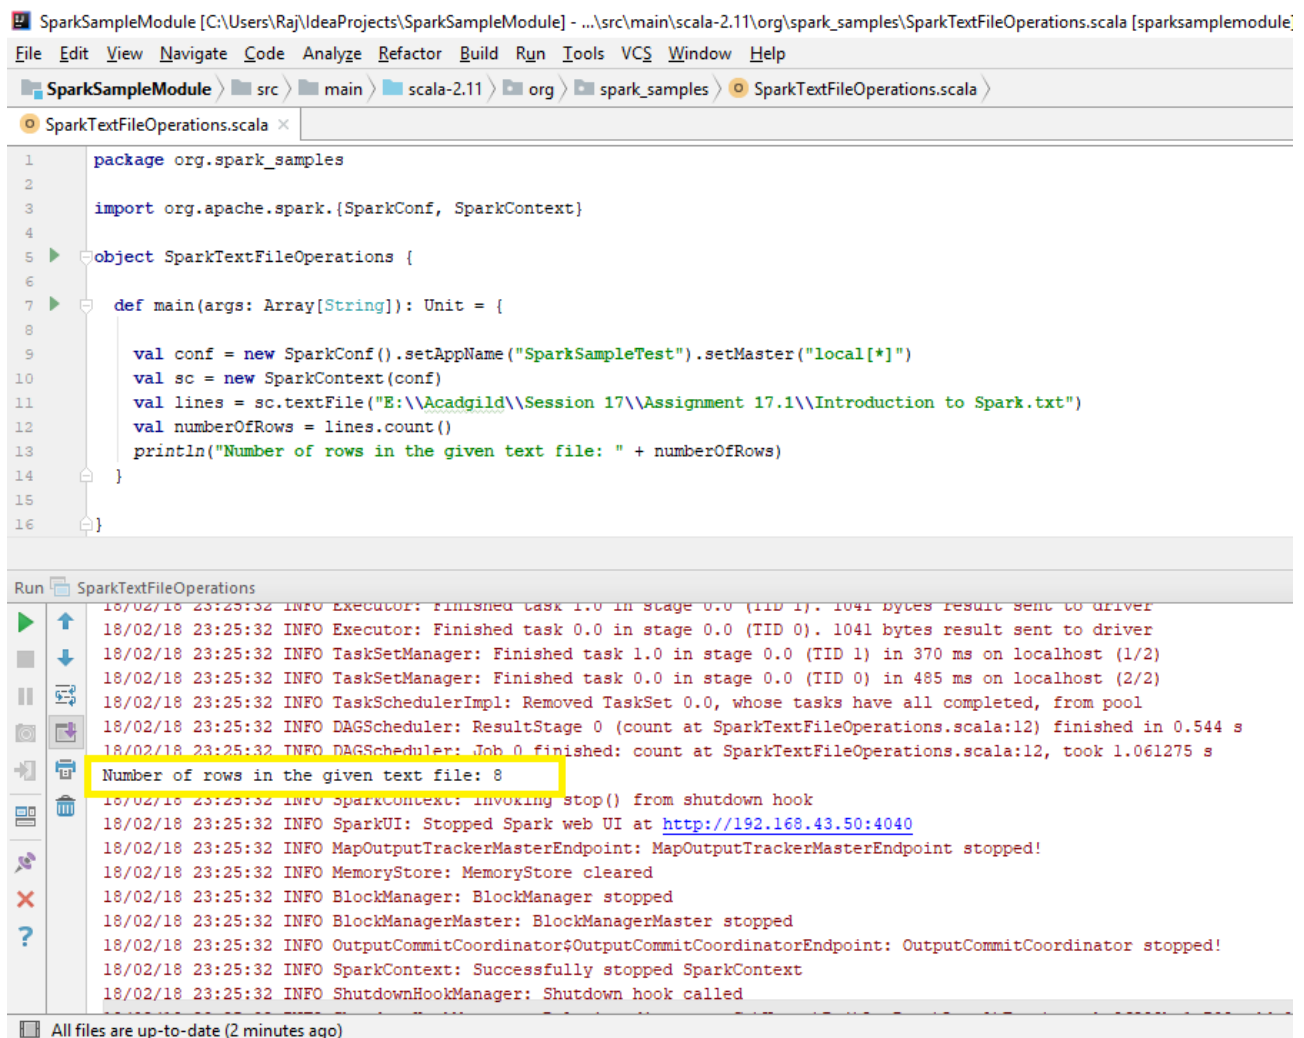


## Task 1

1. Write a program to read a text file and print the number of rows of data in the document.

```
object SparkTextFileOperations {  
  def main(args: Array[String]): Unit = {  
    val conf = new SparkConf().setAppName("SparkTest").setMaster("local[*]")  
    val sc = new SparkContext(conf)  
    val lines = sc.textFile("D:\\Acadgild\\Session 19\\Assignment 19\\Introduction to Spark.txt")  
    val numberOfRows = lines.count()  
    println("Number of rows in the given text file: " + numberOfRows)  
  }  
}
```



The screenshot shows an IDE window for a project named 'SparkSampleModule'. The file 'SparkTextFileOperations.scala' is open, showing the following code:

```
1 package org.spark_samples  
2  
3 import org.apache.spark.{SparkConf, SparkContext}  
4  
5 object SparkTextFileOperations {  
6  
7   def main(args: Array[String]): Unit = {  
8  
9     val conf = new SparkConf().setAppName("SparkSampleTest").setMaster("local[*]")  
10    val sc = new SparkContext(conf)  
11    val lines = sc.textFile("E:\\Acadgild\\Session 17\\Assignment 17.1\\Introduction to Spark.txt")  
12    val numberOfRows = lines.count()  
13    println("Number of rows in the given text file: " + numberOfRows)  
14  }  
15  
16 }
```

The Run console shows the following output:

```
18/02/18 23:25:32 INFO Executor: Finished task 1.0 in stage 0.0 (TID 1). 1041 bytes result sent to driver  
18/02/18 23:25:32 INFO Executor: Finished task 0.0 in stage 0.0 (TID 0). 1041 bytes result sent to driver  
18/02/18 23:25:32 INFO TaskSetManager: Finished task 1.0 in stage 0.0 (TID 1) in 370 ms on localhost (1/2)  
18/02/18 23:25:32 INFO TaskSetManager: Finished task 0.0 in stage 0.0 (TID 0) in 485 ms on localhost (2/2)  
18/02/18 23:25:32 INFO TaskSchedulerImpl: Removed TaskSet 0.0, whose tasks have all completed, from pool  
18/02/18 23:25:32 INFO DAGScheduler: ResultStage 0 (count at SparkTextFileOperations.scala:12) finished in 0.544 s  
18/02/18 23:25:32 INFO DAGScheduler: Job 0 finished: count at SparkTextFileOperations.scala:12, took 1.061275 s  
Number of rows in the given text file: 8  
18/02/18 23:25:32 INFO SparkContext: Invoking stop() from shutdown hook  
18/02/18 23:25:32 INFO SparkUI: Stopped Spark web UI at http://192.168.43.50:4040  
18/02/18 23:25:32 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!  
18/02/18 23:25:32 INFO MemoryStore: MemoryStore cleared  
18/02/18 23:25:32 INFO BlockManager: BlockManager stopped  
18/02/18 23:25:32 INFO BlockManagerMaster: BlockManagerMaster stopped  
18/02/18 23:25:32 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!  
18/02/18 23:25:32 INFO SparkContext: Successfully stopped SparkContext  
18/02/18 23:25:32 INFO ShutdownHookManager: Shutdown hook called
```

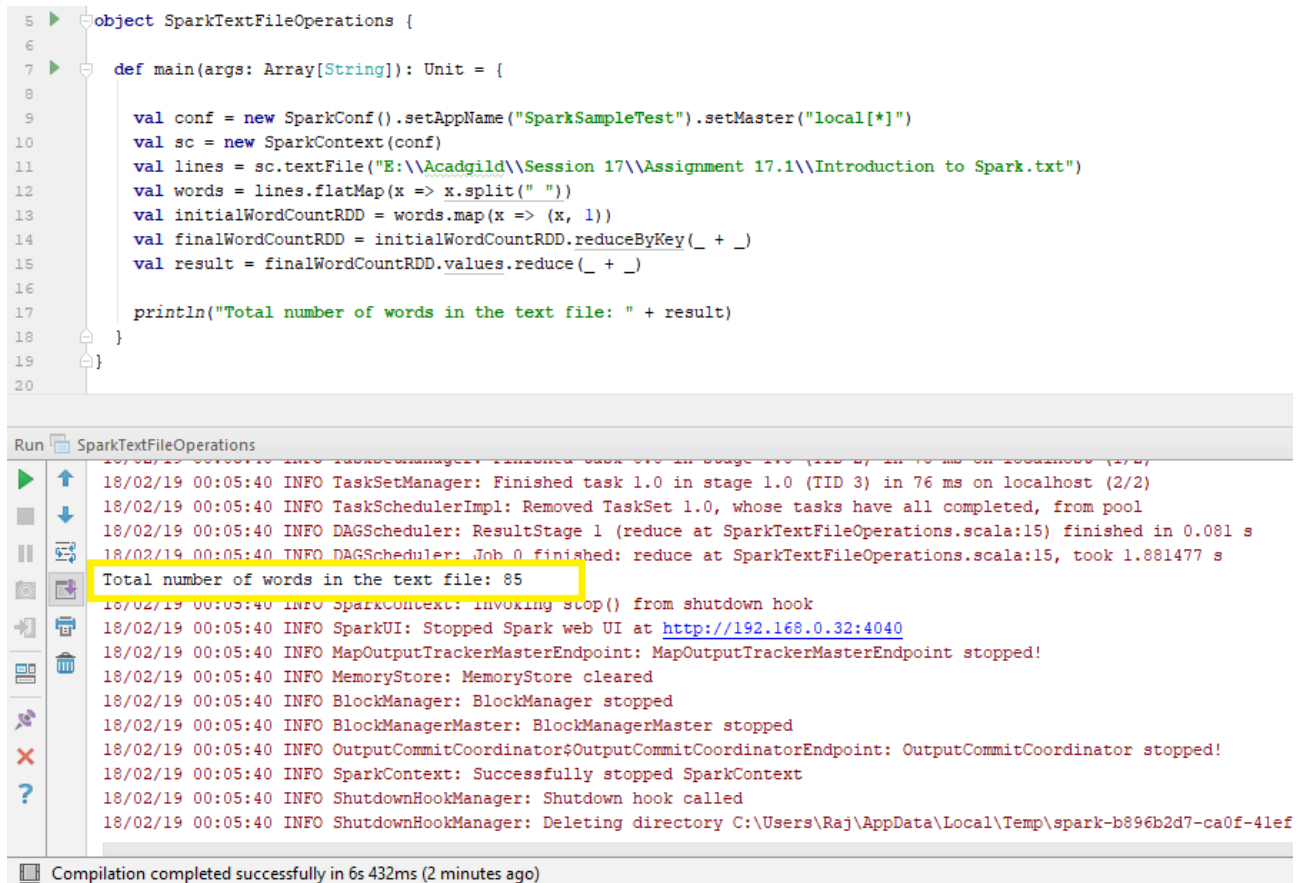
At the bottom, it says 'All files are up-to-date (2 minutes ago)'.

Output:

Number of rows in the given text file: 8

2. Write a program to read a text file and print the number of words in the document.

```
val lines = sc.textFile("D:\\Acadgild\\Session 19\\Assignment 19\\Introduction to Spark.txt")
val words = lines.flatMap(x => x.split(" "))
val initialWordCountRDD = words.map(x => (x, 1))
val finalWordCountRDD = initialWordCountRDD.reduceByKey(_ + _)
val result = finalWordCountRDD.values.reduce(_ + _)
println("Total number of words in the text file: " + result)
```



```
5 object SparkTextFileOperations {
6
7 def main(args: Array[String]): Unit = {
8
9     val conf = new SparkConf().setAppName("SparkSampleTest").setMaster("local[*]")
10    val sc = new SparkContext(conf)
11    val lines = sc.textFile("E:\\Acadgild\\Session 17\\Assignment 17.1\\Introduction to Spark.txt")
12    val words = lines.flatMap(x => x.split(" "))
13    val initialWordCountRDD = words.map(x => (x, 1))
14    val finalWordCountRDD = initialWordCountRDD.reduceByKey(_ + _)
15    val result = finalWordCountRDD.values.reduce(_ + _)
16
17    println("Total number of words in the text file: " + result)
18  }
19 }
20
```

Run SparkTextFileOperations

```
18/02/19 00:05:40 INFO TaskSetManager: Finished task 1.0 in stage 1.0 (TID 3) in 76 ms on localhost (2/2)
18/02/19 00:05:40 INFO TaskSchedulerImpl: Removed TaskSet 1.0, whose tasks have all completed, from pool
18/02/19 00:05:40 INFO DAGScheduler: ResultStage 1 (reduce at SparkTextFileOperations.scala:15) finished in 0.081 s
18/02/19 00:05:40 INFO DAGScheduler: Job 0 finished: reduce at SparkTextFileOperations.scala:15, took 1.881477 s
Total number of words in the text file: 85
18/02/19 00:05:40 INFO SparkContext: Invoking stop() from shutdown hook
18/02/19 00:05:40 INFO SparkUI: Stopped Spark web UI at http://192.168.0.32:4040
18/02/19 00:05:40 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
18/02/19 00:05:40 INFO MemoryStore: MemoryStore cleared
18/02/19 00:05:40 INFO BlockManager: BlockManager stopped
18/02/19 00:05:40 INFO BlockManagerMaster: BlockManagerMaster stopped
18/02/19 00:05:40 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
18/02/19 00:05:40 INFO SparkContext: Successfully stopped SparkContext
18/02/19 00:05:40 INFO ShutdownHookManager: Shutdown hook called
18/02/19 00:05:40 INFO ShutdownHookManager: Deleting directory C:\Users\Raj\AppData\Local\Temp\spark-b896b2d7-ca0f-41ef
```

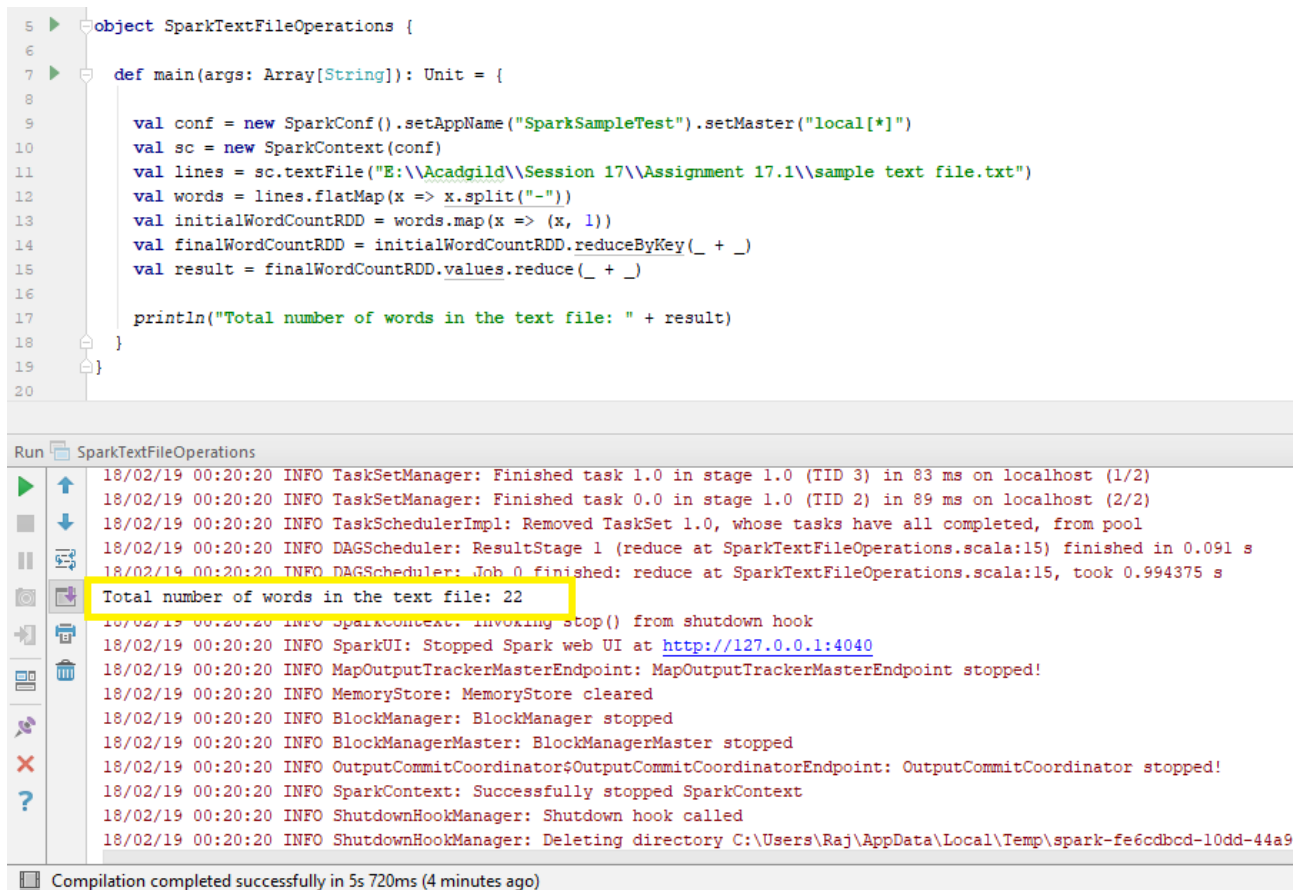
Compilation completed successfully in 6s 432ms (2 minutes ago)

Output:

Total number of words in the text file: 85

3. We have a document where the word separator is -, instead of space. Write a spark code, to obtain the count of the total number of words present in the document.

```
val lines = sc.textFile("D:\\Acadgild\\Session 19\\Assignment 19\\sample text file.txt")
val words = lines.flatMap(x => x.split("-"))
val initialWordCountRDD = words.map(x => (x, 1))
val finalWordCountRDD = initialWordCountRDD.reduceByKey(_ + _)
val result = finalWordCountRDD.values.reduce(_ + _)
println("Total number of words in the text file: " + result)
```



The screenshot shows an IDE with a Scala file named `SparkTextFileOperations`. The code defines a `main` function that reads a text file, splits words by hyphens, and counts them using RDD operations. The output of the program is displayed in the console, showing the total word count of 22. The console also displays various Spark logs, including task completion and context shutdown messages.

```
object SparkTextFileOperations {
  def main(args: Array[String]): Unit = {
    val conf = new SparkConf().setAppName("SparkSampleTest").setMaster("local[*]")
    val sc = new SparkContext(conf)
    val lines = sc.textFile("E:\\Acadgild\\Session 17\\Assignment 17.1\\sample text file.txt")
    val words = lines.flatMap(x => x.split("-"))
    val initialWordCountRDD = words.map(x => (x, 1))
    val finalWordCountRDD = initialWordCountRDD.reduceByKey(_ + _)
    val result = finalWordCountRDD.values.reduce(_ + _)
    println("Total number of words in the text file: " + result)
  }
}
```

Run SparkTextFileOperations

```
18/02/19 00:20:20 INFO TaskSetManager: Finished task 1.0 in stage 1.0 (TID 3) in 83 ms on localhost (1/2)
18/02/19 00:20:20 INFO TaskSetManager: Finished task 0.0 in stage 1.0 (TID 2) in 89 ms on localhost (2/2)
18/02/19 00:20:20 INFO TaskSchedulerImpl: Removed TaskSet 1.0, whose tasks have all completed, from pool
18/02/19 00:20:20 INFO DAGScheduler: ResultStage 1 (reduce at SparkTextFileOperations.scala:15) finished in 0.091 s
18/02/19 00:20:20 INFO DAGScheduler: Job 0 finished: reduce at SparkTextFileOperations.scala:15, took 0.994375 s
Total number of words in the text file: 22
18/02/19 00:20:20 INFO SparkContext: Invoking stop() from shutdown hook
18/02/19 00:20:20 INFO SparkUI: Stopped Spark web UI at http://127.0.0.1:4040
18/02/19 00:20:20 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
18/02/19 00:20:20 INFO MemoryStore: MemoryStore cleared
18/02/19 00:20:20 INFO BlockManager: BlockManager stopped
18/02/19 00:20:20 INFO BlockManagerMaster: BlockManagerMaster stopped
18/02/19 00:20:20 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
18/02/19 00:20:20 INFO SparkContext: Successfully stopped SparkContext
18/02/19 00:20:20 INFO ShutdownHookManager: Shutdown hook called
18/02/19 00:20:20 INFO ShutdownHookManager: Deleting directory C:\Users\Raj\AppData\Local\Temp\spark-fe6cdbc0-10dd-44a9
```

Compilation completed successfully in 5s 720ms (4 minutes ago)

Output:

Total number of words in the text file: 22

## Task 2

### Problem Statement 1:

1. Here is the code snippet I have written in Scala to create a tupled RDD on given data:

```
object Assignment19 {
  def main(args: Array[String]): Unit = {
    val conf = new SparkConf().setAppName("SparkSampleTest").setMaster("local[*]")
    val sc = new SparkContext(conf)
    val lines = sc.textFile("D:\\Acadgild\\Session 19\\Assignment 19\\19_Dataset.txt")
    val tupled_rdd = lines.map(x => {
      val record = x.split(",").toList
      (record.apply(0), record.apply(1), record.apply(2), record.apply(3), record.apply(4))
    })
    tupled_rdd.collect().map(x => println(x._1 + "," + x._2 + "," + x._3))
  }
}
```

Output:

```
18/02/19 22:12:25 INFO DAGScheduler: Job 0 finished: collect at Assignment17_2.scala:16, took 0.626890 s
Mathew,science,grade-3
Mathew,history,grade-2
Mark,maths,grade-2
Mark,science,grade-1
John,history,grade-1
John,maths,grade-2
Lisa,science,grade-1
Lisa,history,grade-3
Andrew,maths,grade-1
Andrew,science,grade-3
Andrew,history,grade-1
Mathew,science,grade-2
Mathew,history,grade-2
Mark,maths,grade-1
Mark,science,grade-2
John,history,grade-1
John,maths,grade-1
Lisa,science,grade-2
Lisa,history,grade-2
Andrew,maths,grade-1
Andrew,science,grade-3
Andrew,history,grade-2
18/02/19 22:12:25 INFO SparkContext: Invoking stop() from shutdown hook
```

2. Spark code in Scala to find the count of total number of rows present:

```
val lines = sc.textFile("D:\\Acadgild\\Session 19\\Assignment 19\\19_Dataset.txt")
println("Total number of rows present in given dataset: " + lines.count())
```

Output:

```
Total number of rows present in given dataset: 22
```

3. Spark code snippet in Scala to find the distinct number of subjects present in the entire school?

```
val lines = sc.textFile("D:\\Acadgild\\Session 19\\Assignment 19\\19_Dataset.txt")
val distinct_subjects_rdd = lines.map(x => x.split(",") (1)).distinct()
println("Total number of distinct subjects " + distinct_subjects_rdd.count())
```

Output:

```
Total number of distinct subjects 3
```

4. Spark code snippet to get count of students in school whose name is Mathew and marks are 55:

```
val lines = sc.textFile("D:\\Acadgild\\Session 19\\Assignment 19\\19_Dataset.txt")
val name_and_marks_rdd = lines.map(x => (x.split(",")(0), x.split(",")(3).toInt))
val filtered_rdd = name_and_marks_rdd.filter(x => x._1 == "Mathew" && x._2 == 55)
println("Total number of rows where student's name is Mathew and marks are 55: " +
filtered_rdd.count())
```

Output:

```
Total number of rows where student's name is Mathew and marks are 55: 2
```

Problem Statement 2:

1. Spark code snippet in Scala to get the count of students per grade in the school:

```
val lines = sc.textFile("D:\\Acadgild\\Session 19\\Assignment 19\\19_Dataset.txt")
val grades_rdd = lines.map(x => x.split(",") (2))
val initial_grades_count_rdd = grades_rdd.map(x => (x, 1))
val final_grades_count = initial_grades_count_rdd.reduceByKey(_ + _)
final_grades_count.foreach(println)
```

Output:

```
(grade-2, 9)
(grade-3, 4)
(grade-1, 9)
```

2. Spark code in Scala to find the average of each student (Note - Mathew is grade-1, is different from Mathew in some other grade!)

```
val lines = sc.textFile("D:\\Acadgild\\Session 19\\Assignment 19\\19_Dataset.txt")
val name_grades_and_marks_rdd = lines.map(x => ((x.split(",") (0), x.split(",") (2)), x.split(",")
(3).toFloat))
val grouped_rdd = name_grades_and_marks_rdd.groupByKey()
```

```
val average_rdd = grouped_rdd.mapValues(x => (x.sum / x.size))
average_marks_rdd.foreach(println)
```

Output:

```
((Lisa,grade-1),24.0)
((Mark,grade-2),17.5)
((Lisa,grade-2),61.0)
((Andrew,grade-2),77.0)
((Mathew,grade-3),45.0)
((John,grade-1),38.666668)
((Andrew,grade-1),43.666668)
((John,grade-2),74.0)
((Lisa,grade-3),86.0)
((Mathew,grade-2),65.666664)
((Mark,grade-1),84.0)
((Andrew,grade-3),35.0)
```

3. Spark code in Scala to get the average score of students in each subject across all grades:

```
val lines = sc.textFile("D:\\Acadgild\\Session 19\\Assignment 19\\19_Dataset.txt")
val name_subject_and_marks_rdd = lines.map(x => ((x.split(",") (0), x.split(",") (1)), x.split(",")
(3).toFloat))
val grouped_rdd = name_subject_and_marks_rdd.groupByKey()
val average_rdd = grouped_rdd.mapValues(x => (x.sum / x.size))
average_marks_rdd.foreach(println)
```

Output:

```
((Lisa,history),92.0)
((Mark,maths),57.5)
((Mark,science),44.0)
((Andrew,science),35.0)
((John,history),40.5)
((Mathew,science),50.0)
((Lisa,science),24.0)
((Andrew,maths),28.5)
((Andrew,history),75.5)
((Mathew,history),71.0)
((John,maths),54.5)
```

4. Spark code in Scala to get the average score of students in each subject per grade:

```
val lines = sc.textFile("D:\\Acadgild\\Session 19\\Assignment 19\\19_Dataset.txt")
val name_subject_grade_and_marks_rdd = lines.map(x => ((x.split(",") (0), x.split(",") (1),
x.split(",") (2)), x.split(",") (3).toFloat))
val grouped_rdd = name_subject_grade_and_marks_rdd.groupByKey()
val average_rdd = grouped_rdd.mapValues(x => (x.sum / x.size))
average_marks_rdd.foreach(println)
```

Output:

```
((Mark,maths,grade-2),23.0)
((Lisa,history,grade-3),86.0)
((Andrew,science,grade-3),35.0)
((Mark,science,grade-2),12.0)
((Mathew,history,grade-2),71.0)
((Andrew,history,grade-1),74.0)
((John,history,grade-1),40.5)
((John,maths,grade-1),35.0)
((Andrew,history,grade-2),77.0)
((John,maths,grade-2),74.0)
((Andrew,maths,grade-1),28.5)
((Mathew,science,grade-3),45.0)
((Mark,maths,grade-1),92.0)
((Mark,science,grade-1),76.0)
((Mathew,science,grade-2),55.0)
((Lisa,science,grade-2),24.0)
((Lisa,history,grade-2),98.0)
((Lisa,science,grade-1),24.0)
```

5. For all students in grade-2, how many have average score greater than 50?

```
val lines = sc.textFile("D:\\Acadgild\\Session 19\\Assignment 19\\19_Dataset.txt")
val grade_2_rdd = lines.filter(x => (x.split(",")(2) == "grade-2"))
val name_and_marks_rdd = grade_2_rdd.map(x => (x.split(",")(0), x.split(",")(3).toFloat))
val grouped_rdd = name_and_marks_rdd.groupByKey()
val average_rdd = grouped_rdd.mapValues(x => (x.sum / x.size))
val average_greater_than_50_rdd = average_rdd.filter(x => (x._2 > 50.0))
println("Total number of students who got average of more than 50 in grade-2: " +
average_greater_than_50_rdd.count())
```

Output:

```
Total number of students who got average of more than 50 in grade-2: 4
```

Problem Statement 3:

Are there any students in the college that satisfy the below criteria:

1. Average score per student\_name across all grades is same as average score per student\_name per grade.

```
val lines = sc.textFile("D:\\Acadgild\\Session 19\\Assignment 19\\19_Dataset.txt")
val name_and_marks_rdd = lines.map(x => (x.split(",") (0), x.split(",") (3).toFloat))
val grouped_rdd = name_and_marks_rdd.groupByKey()
val average_rdd = grouped_rdd.mapValues(x => (x.sum / x.size))
average_rdd.foreach(println)
val name_grade_and_marks_rdd = lines.map(x => ((x.split(",") (0), x.split(",") (2)), x.split(",") (3).toFloat))
val grouped_rdd2 = name_grade_and_marks_rdd.groupByKey()
val average_rdd2 = grouped_rdd2.mapValues(x => (x.sum / x.size))
val simplified_average_rdd2 = average_rdd2.map(x => (x._1._1, x._2))
```

```
simplified_average_rdd2.foreach(println)
val res = average_rdd.intersection(simplified_average_rdd2)
println("Number of students who satisfy the given criteria: " + result.count())
```

```
val lines = sc.textFile("E:\\Acadgild\\Session 17\\Assignment 17.2\\17.2_Dataset.txt")
val name_and_marks_rdd = lines.map(x => (x.split(",")(0), x.split(",")(3).toFloat))
val grouped_rdd = name_and_marks_rdd.groupByKey()
val average_rdd = grouped_rdd.mapValues(x => (x.sum / x.size))
average_rdd.foreach(println)

val name_grade_and_marks_rdd = lines.map(x => ((x.split(",")(0), x.split(",")(2)), x.split(",")(3).toFloat))
val grouped_rdd2 = name_grade_and_marks_rdd.groupByKey()
val average_rdd2 = grouped_rdd2.mapValues(x => (x.sum / x.size))

val simplified_average_rdd2 = average_rdd2.map(x => (x._1._1, x._2))
simplified_average_rdd2.foreach(println)

val result = average_rdd.intersection(simplified_average_rdd2)
println("Number of students who satisfy the given criteria: " + result.count())
```

Output for average score per student\_name across all grades:

```
(Mark,50.75)
(Mathew,60.5)
(Andrew,46.333332)
(John,47.5)
(Lisa,58.0)
```

Output for average score per student\_name per grade:

```
(Lisa,24.0)
(Andrew,77.0)
(John,38.666668)
(John,74.0)
(Mathew,65.666664)
(Mark,17.5)
(Lisa,61.0)
(Mathew,45.0)
(Andrew,43.666668)
(Lisa,86.0)
(Mark,84.0)
(Andrew,35.0)
```

Final Result:

```
Number of students who satisfy the given criteria: 0
```

Hence, we can conclude that there are no students in the college who satisfy the given criteria of falling under both of these results on average scores.